

Process of Building an AI Model for NLP

Building an AI model for NLP involves several key steps, from data preparation to model deployment.

1. Define the Problem

- Clearly define the goal (classification, translation, chatbot, etc.).
- Example: 'Build a model to classify emails as spam or not spam.'

2. Collect and Preprocess Data

- Collect data from web scraping, APIs, public datasets.
- Preprocessing steps: tokenisation, stopwords removal, stemming, lemmatization.

3. Convert Text into Numerical Form

- Traditional: Bag of Words (BoW), TF-IDF.
- Advanced: Word2Vec, GloVe, BERT embeddings.

4. Choose the Right Model

- Machine Learning: Naïve Bayes, SVM, Logistic Regression.
- Deep Learning: LSTMs, GRUs, Transformers (BERT, GPT-4).

5. Train & Evaluate the Model

- Train using labeled data.
- Evaluate with accuracy, F1-score, precision, recall.

6. Fine-Tune & Optimise

- Adjust hyperparameters, use dropout & regularisation.
- Transfer learning with pre-trained models (BERT, GPT).

7. Deploy & Monitor the Model

- Convert to API (Flask, FastAPI).
- Deploy on AWS, Google Cloud, Hugging Face.

Summary of the NLP Model Pipeline:

1. Define Problem - Select NLP task (classification, translation).
2. Collect Data - Scrape, download, or use APIs.

3. Preprocess Data - Clean text, remove stopwords, tokenize.
4. Vectorize Data - BoW, TF-IDF, embeddings (Word2Vec, BERT).
5. Choose Model - ML (Logistic Regression, SVM) or DL (LSTM, Transformers).
6. Train & Evaluate - Train model, validate, tune hyperparameters.
7. Deploy Model - Convert into API, deploy to cloud.

For simple NLP tasks, use traditional ML models.

For deep NLP understanding, use LSTMs, GRUs, or Transformers (BERT, GPT-4).

For scalable AI applications, fine-tune pre-trained Transformer models.